

Type *Markdown* and LaTeX:  $\alpha^2$

```
In [1]: #import libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: #import dataset
df=pd.read_csv(r"E:\154\drive-download-20230731T110444Z-001\20_states.csv").dropna
df
```

Out[2]:

	id	name	country_id	country_code	country_name	state_code	type	latitude	lo
170	3656	Buenos Aires	11	AR	Argentina	B	province	-37.201729	-59
171	3647	Catamarca	11	AR	Argentina	K	province	-28.471588	-65
172	3640	Chaco	11	AR	Argentina	H	province	-27.425718	-59
173	3651	Chubut	11	AR	Argentina	U	province	-43.293425	-65
174	4880	Ciudad Autónoma de Buenos Aires	11	AR	Argentina	C	city	-34.603684	-58
...	...	...	...	...	...	...	...	...	...
4968	2041	Yaracuy	239	VE	Venezuela	U	state	10.339389	-68
4969	2042	Zulia	239	VE	Venezuela	V	state	10.291024	-72
5033	5074	Saint Croix	242	VI	Virgin Islands (US)	SC	district	17.729352	-64
5034	5073	Saint John	242	VI	Virgin Islands (US)	SJ	district	18.335617	-64
5035	5072	Saint Thomas	242	VI	Virgin Islands (US)	ST	district	18.342846	-65

1580 rows × 9 columns



In [3]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1580 entries, 170 to 5035
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id               1580 non-null   int64
1   name             1580 non-null   object
2   country_id       1580 non-null   int64
3   country_code     1580 non-null   object
4   country_name     1580 non-null   object
5   state_code       1580 non-null   object
6   type             1580 non-null   object
7   latitude         1580 non-null   float64
8   longitude        1580 non-null   float64
dtypes: float64(2), int64(2), object(5)
memory usage: 123.4+ KB
```

In [4]: *#to display top 5 rows*  
df.head()

Out[4]:

	id	name	country_id	country_code	country_name	state_code	type	latitude	lon
170	3656	Buenos Aires	11	AR	Argentina	B	province	-37.201729	-59.8
171	3647	Catamarca	11	AR	Argentina	K	province	-28.471588	-65.7
172	3640	Chaco	11	AR	Argentina	H	province	-27.425718	-59.0
173	3651	Chubut	11	AR	Argentina	U	province	-43.293425	-65.
174	4880	Ciudad Autónoma de Buenos Aires	11	AR	Argentina	C	city	-34.603684	-58.5

## Data cleaning and Pre-Processing

In [5]: *#To find null values*  
df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1580 entries, 170 to 5035
Data columns (total 9 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   id              1580 non-null   int64
 1   name            1580 non-null   object
 2   country_id      1580 non-null   int64
 3   country_code    1580 non-null   object
 4   country_name    1580 non-null   object
 5   state_code      1580 non-null   object
 6   type            1580 non-null   object
 7   latitude        1580 non-null   float64
 8   longitude       1580 non-null   float64
dtypes: float64(2), int64(2), object(5)
memory usage: 123.4+ KB
```

In [6]: *# To display summary of statistics*  
df.describe()

Out[6]:

	id	country_id	latitude	longitude
<b>count</b>	1580.000000	1580.000000	1580.000000	1580.000000
<b>mean</b>	2685.916456	134.000633	26.988930	15.009671
<b>std</b>	1611.169440	66.055166	19.635279	66.200355
<b>min</b>	48.000000	11.000000	-54.805400	-178.116500
<b>25%</b>	1339.750000	75.000000	13.752013	-7.622388
<b>50%</b>	2210.500000	139.000000	30.887089	11.665277
<b>75%</b>	4013.250000	178.000000	42.938004	45.682217
<b>max</b>	5220.000000	242.000000	77.874972	166.649935

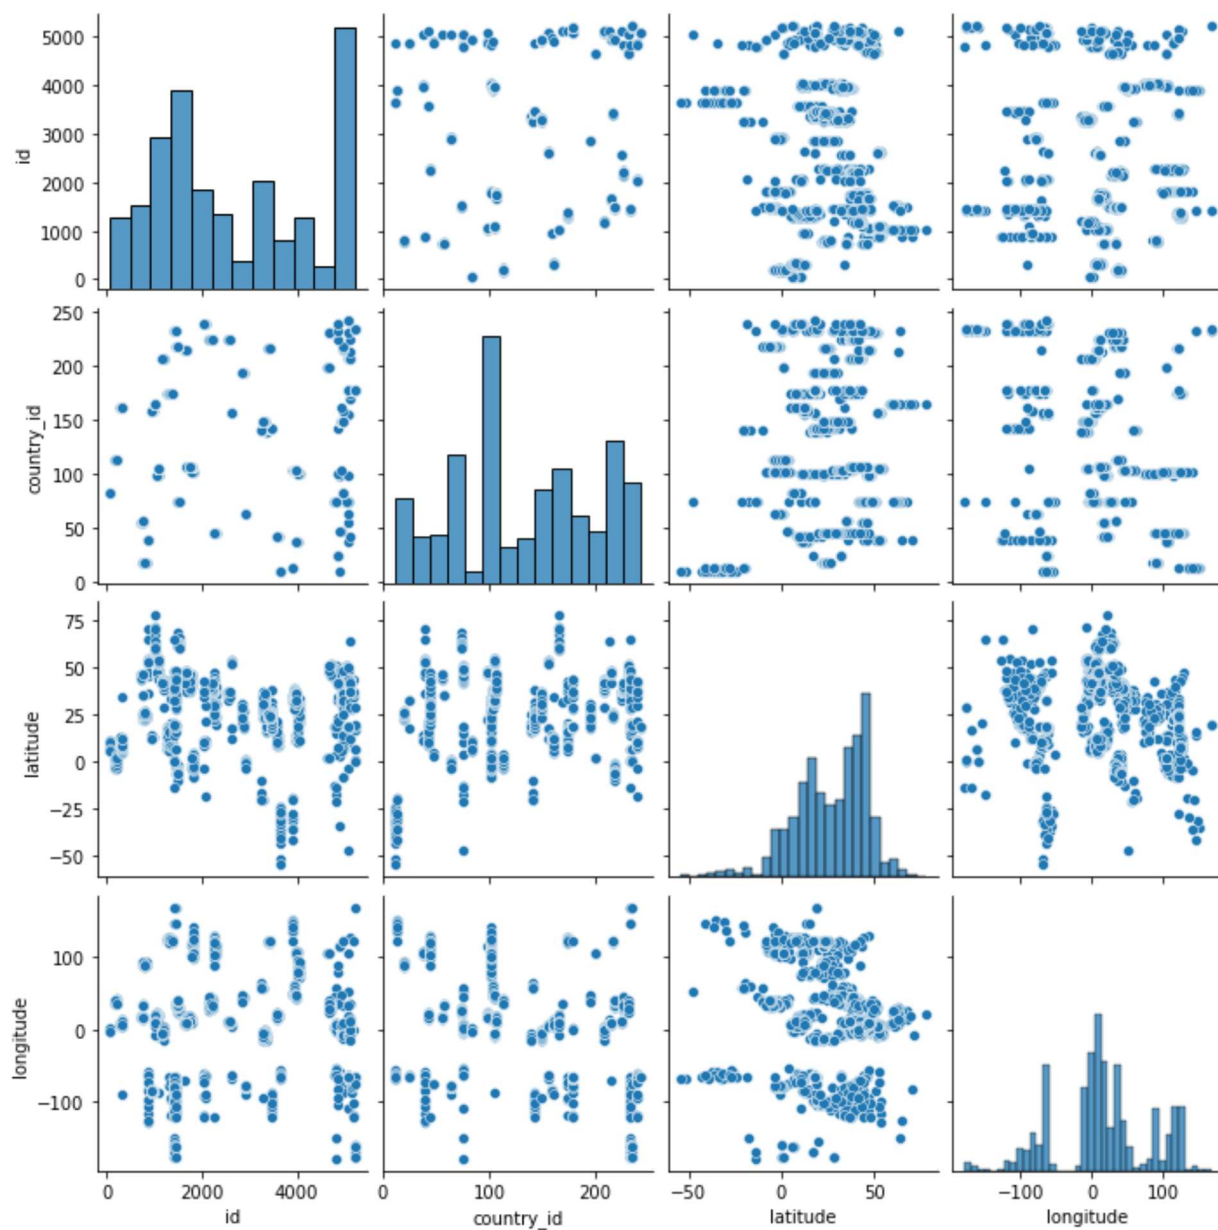
In [7]: *#To Display column heading*  
df.columns

Out[7]: Index(['id', 'name', 'country\_id', 'country\_code', 'country\_name',  
          'state\_code', 'type', 'latitude', 'longitude'],  
          dtype='object')

## EDA and VISUALIZATION

```
In [8]: sns.pairplot(df)
```


```
Out[8]: <seaborn.axisgrid.PairGrid at 0x159f74529d0>
```



```
In [9]: sns.distplot(df['longitude'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).  
warnings.warn(msg, FutureWarning)

```
Out[9]: <AxesSubplot:xlabel='longitude', ylabel='Density'>
```




```
In [10]: df1=df[['id','country_id', 'latitude','longitude']]
```

## Plot Using Heat Map

```
In [11]: sns.heatmap(df1.corr())
```

```
Out[11]: <AxesSubplot:>
```



## To Train The Model-Model Building

we are going to train Linera Regression Model;We need to split out data into two variables x and y where x is independent variable(input) and y is dependent on x(output) we could ignore address column as it required for our model

```
In [12]: x=df1[['id','country_id', 'latitude']]
y=df1[ 'longitude']
```

## To Split my dataset into training and test data

```
In [13]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [14]: from sklearn.linear_model import LinearRegression
lr= LinearRegression()
lr.fit(x_train,y_train)
```

Out[14]: LinearRegression()

```
In [15]: lr.intercept_
```

Out[15]: 69.88512017384633


```
In [16]: coeff = pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

Out[16]:

	Co-efficient
id	-0.007278
country_id	-0.217729
latitude	-0.236816

```
In [17]: prediction = lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[17]: <matplotlib.collections.PathCollection at 0x159f9042640>



```
In [18]: lr.score(x_test,y_test)
```

```
Out[18]: 0.10055703466560706
```

## Accuracy

```
In [19]: lr.score(x_test,y_test)
```

```
Out[19]: 0.10055703466560706
```

```
In [20]: lr.score(x_train,y_train)
```

```
Out[20]: 0.0904811044679853
```

```
In [21]: from sklearn.linear_model import Ridge,Lasso
```

```
In [22]: rr=Ridge(alpha=10)
         rr.fit(x_train,y_train)
```

```
Out[22]: Ridge(alpha=10)
```

```
In [23]: rr.score(x_test,y_test)
```

```
Out[23]: 0.10055692068940869
```

```
In [24]: la =Lasso(alpha=10)
         la.fit(x_train,y_train)
```

```
Out[24]: Lasso(alpha=10)
```

```
In [25]: la.score(x_test,y_test)
```

```
Out[25]: 0.09996733626958476
```

## ElasticNet

```
In [26]: from sklearn.linear_model import ElasticNet
         en = ElasticNet()
         en.fit(x_train,y_train)
```

```
Out[26]: ElasticNet()
```

```
In [27]: print(en.coef_)

[-0.00727859 -0.21763718 -0.23528058]
```

```
In [28]: print(en.intercept_)

69.83122292494286
```

In [29]: `print(en.predict(x_test))`

```
[-4.22485152e+00  1.12566122e+01  3.52311015e+01 -6.02626880e+00
 6.50846781e+00  1.17778772e+01  4.32514695e+01  6.51429911e+00
-2.09029462e+01  2.37112117e+01 -1.09340160e+01  3.42660808e+01
-4.32407664e+00  1.25859531e-01  9.63913303e+00  2.11002486e+01
 4.35922537e+00  1.04490732e+01  4.37809134e+01  2.43111730e+01
 2.03274788e+01 -6.21787685e+00  2.31463592e+01  2.65334598e+01
 2.63910542e+01  1.97665144e+01  5.44014220e+01  2.93340491e+01
 6.42631596e+00  3.00716425e+01  1.33007572e+01  2.16071254e+00
 2.36315904e+01  2.38769891e+01 -2.42455387e+01  1.98402507e+01
 3.38810100e+00  2.66045166e+01  3.60171637e+01  2.97632254e+01
 2.06525774e+01  5.42817474e+01  2.08586698e+01 -2.09721671e+01
 3.51532005e+01  2.55216500e+01 -7.17057998e-02  2.02255619e+01
-1.11044568e+01  4.91257303e+00  1.54456160e+00  2.44017844e+01
 1.07826037e+01  6.55170807e+00  1.97115516e+01  7.10066803e+00
-8.09098789e+00 -6.19920185e-01  1.36962229e+01  1.88686876e+01
 3.09161514e+01 -7.88755020e+00 -1.92278229e+01  2.95651494e+01
 1.08308905e+01  5.96914035e+00  3.00187408e+01  2.76914353e+01
 9.31109301e+00  7.66158883e+00  2.34114130e+01  5.10491971e+01
 6.56717229e+00  5.43499430e+00  1.18291642e+01 -4.87125718e+00
 4.33333333e+01  3.61000000e+01  7.15000000e+00  1.07100000e+01]
```

In [30]: `print(en.score(x_test,y_test))`

0.10052365459337198

## Evaluation Metrics

In [31]: `from sklearn import metrics`

In [32]: `print("Mean Absolute Error",metrics.mean_absolute_error(y_test,prediction))`

Mean Absolute Error 49.979246514739344

In [33]: `print("Mean Squared Error:",metrics.mean_squared_error(y_test,prediction))`

Mean Squared Error: 4195.473860876993

In [34]: `print("Root Mean Absolute Error:",np.sqrt(metrics.mean_squared_error(y_test,prediction)))`

Root Mean Absolute Error: 64.77247764966995

In [35]: `print("Root Mean Absolute Error:",np.sqrt(metrics.mean_squared_error(y_test,prediction)))`

Root Mean Absolute Error: 64.77247764966995



## Model Saving

```
In [36]: import pickle
```

```
In [37]: filename="prediction"
pickle.dump(lr,open(filename,'wb'))
```

```
In [38]: model=pickle.load(open(filename,'rb'))
```

```
In [39]: real=[[10,20,30],[11,45,10]]
```

```
In [40]: result =model.predict(real)
```

```
In [41]: print(result)

[58.35327921 57.63907598]
```

```
In [ ]:
```